

Predicting Atlanta Falcons NFL Receiver Touchdowns with Regression Modelling

Bhargav Ashok

2025-05-04

Installing packages

Loading our Atlanta Falcons Data

```
Atlanta_Falcons_data <- read_excel("Atlanta_Falcons_data.xlsx")
head(Atlanta_Falcons_data)
```

```
## # A tibble: 6 x 18
##   Rk Player      From To      G Pos      AV  Tgt  Rec 'Ctch%'  Yds 'Y/R'
##   <dbl> <chr>      <dbl> <dbl> <dbl> <chr> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1     1 Julio Jon~ 2011 2020 135 WR      119 1320 848 0.642 12896 15.2
## 2     2 Roddy Whi~ 2005 2015 171 WR      107 1377 808 0.587 10863 13.4
## 3     3 Terance M~ 1994 2001 126 WR       67 989 573 0.579 7349 12.8
## 4     4 Alfred Je~ 1975 1983 110 WR       61  NA 360  NA    6267 17.4
## 5     5 Andre Ris~ 1990 1994 78  WR       53 463 423  NA    5633 13.3
## 6     6 Jim Mitch~ 1969 1979 155 TE       47  NA 305  NA    4358 14.3
## # i 6 more variables: TD <dbl>, Lng <dbl>, 'Y/Tgt' <dbl>, 'R/G' <dbl>,
## #   'Y/G' <dbl>, Fmb <dbl>
```

Preparing our data, cleaning and training

Here, we are cleaning our data with the tidyr package and then splitting our data into a trained and tested set.

```
vars_needed <- c("TD", "Tgt", "Rec", "Ctch%", "Yds")
clean_df    <- Atlanta_Falcons_data %>% drop_na(all_of(vars_needed))

set.seed(123)
n          <- nrow(clean_df)
train_i    <- sample(n, 0.8*n)
train_df   <- clean_df[ train_i, ]
test_df    <- clean_df[-train_i, ]
```

fitting our Model (regular MLR)

Next, we fit our model using the equation: $Y = (\beta)_0 + (\beta)_1X_1 + (\beta)_2X_2 + \dots + (\beta)_nX_n + \text{error}$ (assuming normal distribution)

```
mlr_fit <- lm(TD ~ Tgt + Rec + `Ctch%` + Yds, data = train_df)
summary(mlr_fit)
```

```
##
## Call:
## lm(formula = TD ~ Tgt + Rec + `Ctch%` + Yds, data = train_df)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -10.4458  -0.7242  -0.1969   0.7605  10.6716
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -0.3606742  0.7621684  -0.473   0.6368
## Tgt          0.0312459  0.0135714   2.302   0.0228 *
## Rec         -0.0044023  0.0176458  -0.249   0.8034
## `Ctch%`      0.5479172  1.0636249   0.515   0.6073
## Yds          0.0020808  0.0008571   2.428   0.0165 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 2.423 on 138 degrees of freedom
## Multiple R-squared:  0.9028, Adjusted R-squared:  0.8999
## F-statistic: 320.3 on 4 and 138 DF,  p-value: < 2.2e-16
```


The summary statistics are displayed above. Note the coefficients for the MLR equation and the adjusted R-squared Value.

MLR fit data display (looking at model data)

Then, we display plots for our fitted data


```
hist(resid(mlr_fit), main="OLS Train Residuals", xlab="Residual")
```

OLS Train Residuals



```
qqnorm(resid(mlr_fit)); qqline(resid(mlr_fit))
```

Normal Q-Q Plot



```
vif(mlr_fit)
```


```
##      Tgt      Rec    'Ctch%'      Yds  
## 110.938423  76.817866  1.126715  32.696953
```

test data display (looking at unseen test data)

After that, we look at our unseen test data plots


```
test_df$pred_ols <- predict(mlr_fit, newdata = test_df)  
resid_test_ols  <- test_df$TD - test_df$pred_ols  
  
# Plots  
hist(resid_test_ols, main="OLS Test Residuals", xlab="Residual")
```

OLS Test Residuals



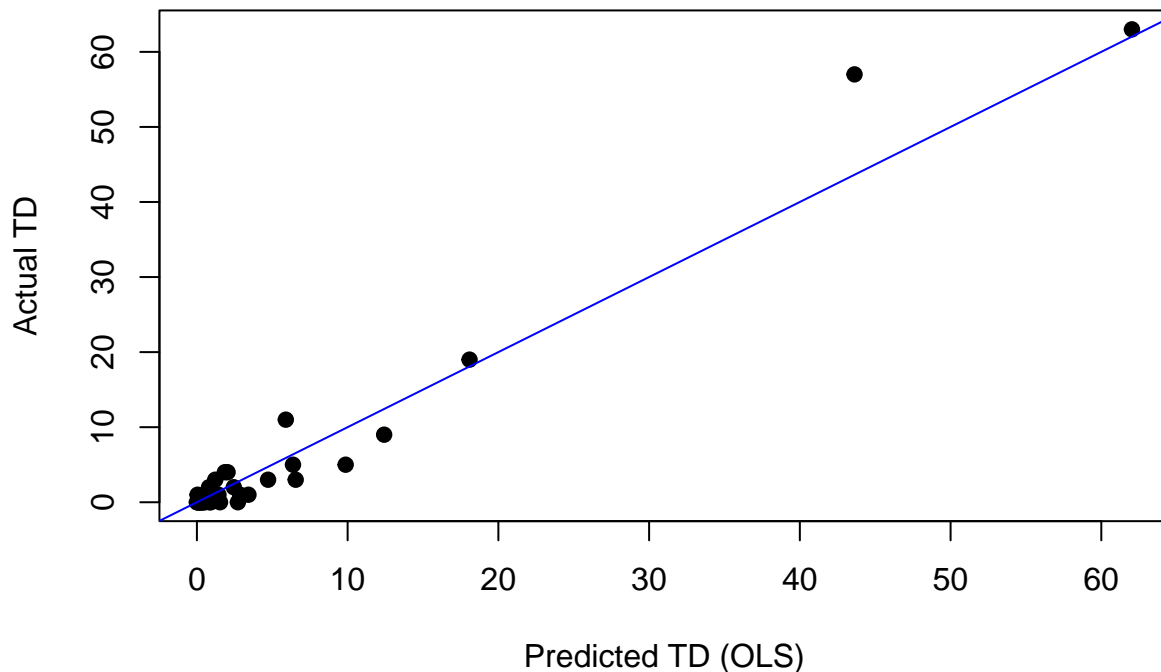
```
qqnorm(resid_test_ols); qqline(resid_test_ols)
```

Normal Q-Q Plot



```
plot(
  test_df$pred_ols, test_df$TD,
  xlab="Predicted TD (OLS)", ylab="Actual TD",
  main="Test: OLS Pred vs Actual", pch=19
)
abline(0,1,col="blue")
```

Test: OLS Pred vs Actual



We see the linear regression line on our predicted test data above.

Using Cross validation + Ridge/Lasso Regression

After looking at our model and our VIF (Variance Inflation Factor, which shows how standardized our data is in terms of multicollinearity), we see that the $VIF > 10$, which indicates high multicollinearity. This is not ideal for this scenario because the variables need to be standardized to account for the number inflation in multicollinearity. Due to this, we need to switch to a new type of regression model for even more accurate results. We use a technique called K-folds cross validation, where the data is split into multiple subsets and is iterated more than once in order to account for the multicollinearity inflation which is indicated above, as well improving the model to see how accurately it can predict unseen data points. We use the glmnet package for this.

```
# Prepare matrices
x_train <- model.matrix(TD~Tgt+Rec+`Ctch%`+Yds, train_df)[,-1]
y_train <- train_df$TD
x_test  <- model.matrix(TD~Tgt+Rec+`Ctch%`+Yds, test_df)[,-1]
y_test  <- test_df$TD

# Ridge
cv_ridge <- cv.glmnet(x_train,y_train,alpha=0)
best_ridge<- cv_ridge$lambda.min
ridge_mod <- glmnet(x_train,y_train,alpha=0,lambda=best_ridge)
test_df$pred_ridge <- as.numeric(predict(ridge_mod,x_test))

# Lasso
```

```

cv_lasso <- cv.glmnet(x_train,y_train,alpha=1)
best_lasso<- cv_lasso$lambda.min
lasso_mod <- glmnet(x_train,y_train,alpha=1,lambda=best_lasso)
test_df$pred_lasso <- as.numeric(predict(lasso_mod,x_test))

# RMSE
cat("Ridge RMSE:", rmse(y_test, test_df$pred_lasso), "\n")

```

```
## Ridge RMSE: 3.52247
```

```
cat("Lasso RMSE:", rmse(y_test, test_df$pred_lasso), "\n")
```

```
## Lasso RMSE: 3.975591
```


We do a Cross Validation of Ridge and lasso regression to see which one is more accurate. As we can see, Ridge regression has a lower RMSE which is more accurate for our model, so we will plot the Cross validation curve.

Plotting Ridge regression

```

plot(cv_lasso)
abline(v=log(best_lasso),col="red",lty=2)

```

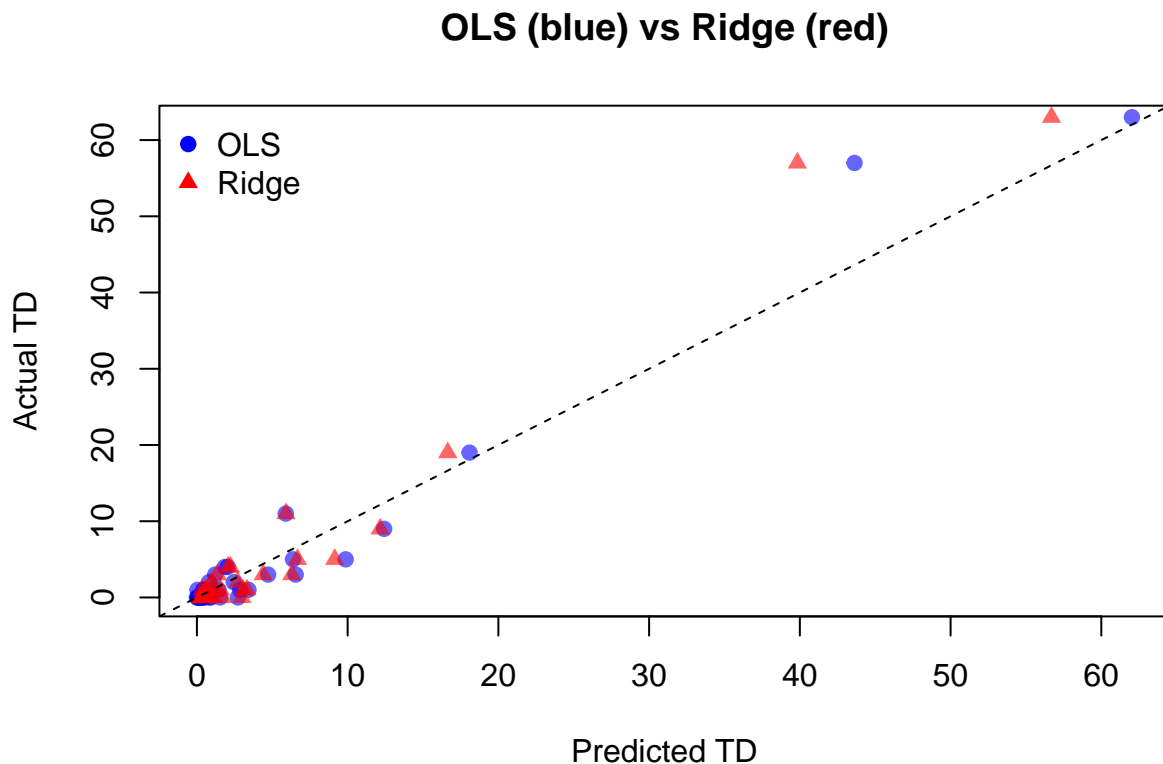


The cross validation ridge plot is shown above. According to “<https://bookdown.org/ssjackson300/Machine-Learning-Lecture-Notes/choosing-lambda.html>”:

“What is plotted is the estimated CV MSE for each value of (log) lambda on the x-axis. The dotted line on the far left indicates the value of lambda which minimizes CV error. The dotted line roughly in the middle of the x-axis indicates the 1-standard-error lambda- recall that this is the maximum value that lambda can take while still falling within the one standard error interval of the minimum-CV lambda. The second line of code has manually added a dot-dash horizontal line at the upper end of the 1-standard deviation interval of the MSE at the minimum-CV lambda to illustrate this point further”. These plots can change with randomization according to our seed number.

Plotting our Comparison graph between MLR and Ridge Regression MLR

```
plot(test_df$pred_ols, test_df$TD,
     xlim = range(c(test_df$pred_ols, test_df$pred_ridge)),
     ylim = range(c(test_df$pred_ols, test_df$pred_ridge)),
     xlab="Predicted TD", ylab="Actual TD",
     main="OLS (blue) vs Ridge (red)", pch=19, col=rgb(0,0,1,0.6))
points(test_df$pred_ridge, test_df$TD, pch=17, col=rgb(1,0,0,0.6))
abline(0,1,lty=2)
legend("topleft", legend=c("OLS", "Ridge"), pch=c(19,17),
      col=c("blue", "red"), bty="n")
```



Finally, we can compare our MLR Ordinary Least Squares Regression (Linear Regression) Model with our Cross-Validated, Ridge Regression Model visually.