

# Predicting Atlanta Falcons NFL Touchdowns with Regression Modelling

Bhargav Ashok

2025-05-16

## 0. Setup - Install Packages: tidyr, readxl, dplyr, car, glmnet, Metrics, 3dscatterplot

### Loading our Atlanta Falcons Data

```
Atlanta_Falcons_data <- read_excel("Atlanta_Falcons_data.xlsx")
head(Atlanta_Falcons_data)
```

```
## # A tibble: 6 x 18
##      Rk Player      From   To      G Pos      AV   Tgt   Rec 'Ctch%'   Yds 'Y/R'
##    <dbl> <chr>      <dbl> <dbl> <dbl> <chr> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1     1  Julio Jon~  2011  2020   135 WR     119  1320   848   0.642 12896  15.2
## 2     2  Roddy Whi~  2005  2015   171 WR     107  1377   808   0.587 10863  13.4
## 3     3  Terance M~  1994  2001   126 WR      67   989   573   0.579  7349  12.8
## 4     4  Alfred Je~  1975  1983   110 WR      61    NA   360   NA     6267  17.4
## 5     5  Andre Ris~  1990  1994    78 WR      53   463   423   NA     5633  13.3
## 6     6  Jim Mitch~  1969  1979   155 TE      47    NA   305   NA     4358  14.3
## # i 6 more variables: TD <dbl>, Lng <dbl>, 'Y/Tgt' <dbl>, 'R/G' <dbl>,
## # 'Y/G' <dbl>, Fmb <dbl>
```

### Preparing our data, cleaning and training

Here, we are cleaning our data with the tidyr package and then splitting our data into a trained and tested set.

```
vars_needed <- c("TD", "Tgt", "Rec", "Ctch%", "Yds")
clean_df    <- Atlanta_Falcons_data %>% drop_na(all_of(vars_needed))

set.seed(123)
n          <- nrow(clean_df)
train_i    <- sample(n, 0.8*n)
train_df   <- clean_df[ train_i, ]
test_df    <- clean_df[-train_i, ]
```

### Fitting our Model (regular MLR)

Next, we fit our model using the equation:  $Y = (\text{Beta})_0 + (\text{Beta})_1X_1 + (\text{Beta})_2X_2 + \dots + (\text{Beta})_nX_n + \text{error}$  (assuming normal distribution)

```
mlr_fit <- lm((TD) ~ Tgt + Rec + `Ctch%` + Yds, data = train_df)
summary(mlr_fit)
```

```
##
## Call:
## lm(formula = (TD) ~ Tgt + Rec + `Ctch%` + Yds, data = train_df)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -10.4458  -0.7242  -0.1969   0.7605  10.6716
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -0.3606742  0.7621684  -0.473   0.6368
## Tgt          0.0312459  0.0135714   2.302   0.0228 *
## Rec         -0.0044023  0.0176458  -0.249   0.8034
## `Ctch%`      0.5479172  1.0636249   0.515   0.6073
## Yds          0.0020808  0.0008571   2.428   0.0165 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 2.423 on 138 degrees of freedom
## Multiple R-squared:  0.9028, Adjusted R-squared:  0.8999
## F-statistic: 320.3 on 4 and 138 DF,  p-value: < 2.2e-16
```

The summary statistics are displayed above. Note the coefficients for the MLR equation and the adjusted R-squared Value.

The MLR equation can be identified and written as: “ $TD = -0.36 + 0.03Tgt - 0.004Rec + 0.55Ctch\% + 0.002Yds$ ”

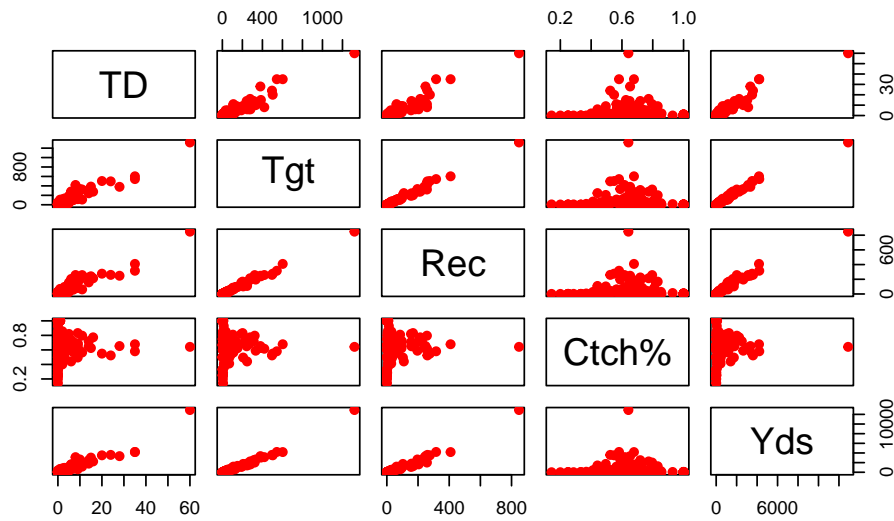
This equation suggests that Catch% has the largest effect on Touchdowns by a factor of 0.548, with yards being a close second with an effect factor of 0.00208. The other factors(reception targets and receptions, meaning how many times the said player was targeted for the throw, and how many times they completed a given catch) seem to have a negative effect, suggesting over fitting. However, this is handled via cross-validation and Ridge regression in the improved MLR model later on. This means that all of these factors influence the number of touchdowns by varying rates based on test statistics performed in this module.

## Data Visualization of the Matrix Scatterplot

Here we see that the scatter plot compares all the different variables together to check for linearity and how each variable influences each other. We may see non-linearity between some variables but that will be accounted for later on.

```
pairs(
  ~ TD + Tgt + Rec + `Ctch%` + Yds,
  data = train_df,
  main = "Matrix Scatterplot of MLR Variables",
  col = "red",
  pch = 19
)
```

## Matrix Scatterplot of MLR Variables




### Data visualization of the 3d scatterplot

Here we are taking the top two linearly correlated variables and plotting them on a 3d scatterplot using the R 3d scatterplot package.

```
scatterplot3d(
  x = train_df$Tgt,
  y = train_df$TD,
  z = train_df$Yds,
  main = "3D Scatterplot: TD ~ Tgt + Yds",
  xlab = "Targets",
  ylab = "Touchdowns",
  zlab = "Yards",
  color = "green",
  pch = 19
)
```

**3D Scatterplot: TD ~ Tgt + Yds**




MLR fit data display (looking at model data)


Then, we display plots for our fitted data (qqnorm and histogram of residuals)

```
hist(resid(mlr_fit), main="OLS Train Residuals", xlab="Residual")
```

**OLS Train Residuals**



```
qqnorm(resid(mlr_fit)); qqline(resid(mlr_fit))
```




From the plots, we see that our trained residual data is roughly normal along with the qqplot being roughly normal as well (with deviation in the extreme values or outliers).

#### test data display (looking at unseen test data)


After that, we look at our unseen test data plots.

```
test_df$pred_ols <- predict(mlr_fit, newdata = test_df)
resid_test_ols   <- test_df$TD - test_df$pred_ols


# Plots
hist(resid_test_ols, main="OLS Test Residuals", xlab="Residual")
```



```
qqnorm(resid_test_ols); qqline(resid_test_ols)
```



```
plot(
  test_df$pred_ols, test_df$TD,
  xlab="Predicted TD (OLS)", ylab="Actual TD",
  main="Test: OLS Pred vs Actual", pch=19
)
abline(0,1,col="blue")
```



We see the linear regression line on our predicted test data above. We don't need to assume normality for our test residuals in a Machine Learning Model (unless a hypothesis test is conducted within a controlled experiment).

### Using Cross validation + Ridge/Lasso Regression

```
vif(mlr_fit)
```

```
##          Tgt          Rec      'Ctch%'      Yds
## 110.938423  76.817866   1.126715  32.696953
```

After looking at our models and our VIF (Variance Inflation Factor, found using car package in R), which shows how standardized our data is in terms of multicollinearity). We see that the  $VIF > 10$ , which indicates high multicollinearity. This is not ideal for this scenario because the variables need to be standardized to account for the number inflation in multicollinearity (since they are supposed to be statistically significant, not reflected in the regular MLR p-value). Due to this, we need to switch to a new type of regression model for even more accurate results. We use a technique called K-folds cross validation, where the data is split into multiple subsets and is iterated more than once in order to account for the multicollinearity inflation which is indicated above, as well improving the model to see how accurately it can predict unseen data points. This involves putting our train/test data into matrices, and then running regularized models called Ridge and Lasso regression respectively. Lasso regression accounts for the absolute value of the important coefficients and shrinks them using a penalty factor. Ridge regression accounts for the squared value of the coefficients and shrinks them with a similar penalty factor (all in the means of regularizing our data (also called hyper parameter tuning)). We use the glmnet package for this.

```
# Prepare matrices
x_train <- model.matrix(TD~Tgt+Rec+`Ctch%`+Yds, train_df)[,-1]
y_train <- train_df$TD
x_test  <- model.matrix(TD~Tgt+Rec+`Ctch%`+Yds, test_df)[,-1]
```

```

y_test <- test_df$TD

# Ridge
cv_ridge <- cv.glmnet(x_train,y_train,alpha=0)
best_ridge<- cv_ridge$lambda.min
ridge_mod <- glmnet(x_train,y_train,alpha=0,lambda=best_ridge)
test_df$pred_ridge <- as.numeric(predict(ridge_mod,x_test))
summary(cv_ridge)

```

```

##           Length Class  Mode
## lambda    100    -none- numeric
## cvm        100    -none- numeric
## cvsd       100    -none- numeric
## cvup       100    -none- numeric
## cvlo       100    -none- numeric
## nzero      100    -none- numeric
## call        4    -none- call
## name        1    -none- character
## glmnet.fit  12    elnet  list
## lambda.min   1    -none- numeric
## lambda.1se   1    -none- numeric
## index       2    -none- numeric

```

```

# Lasso
cv_lasso <- cv.glmnet(x_train,y_train,alpha=1)
best_lasso<- cv_lasso$lambda.min
lasso_mod <- glmnet(x_train,y_train,alpha=1,lambda=best_lasso)
test_df$pred_lasso <- as.numeric(predict(lasso_mod,x_test))
summary(cv_lasso)

```

```

##           Length Class  Mode
## lambda     59    -none- numeric
## cvm         59    -none- numeric
## cvsd        59    -none- numeric
## cvup        59    -none- numeric
## cvlo        59    -none- numeric
## nzero       59    -none- numeric
## call         4    -none- call
## name         1    -none- character
## glmnet.fit  12    elnet  list
## lambda.min   1    -none- numeric
## lambda.1se   1    -none- numeric
## index       2    -none- numeric

```

```

# RMSE
cat("Ridge RMSE:", rmse(y_test, test_df$pred_ridge), "\n")

```

```
## Ridge RMSE: 3.52247
```

```
cat("Lasso RMSE:", rmse(y_test, test_df$pred_lasso), "\n")
```


```
## Lasso RMSE: 3.975591
```



We do a Cross Validation of Ridge and Lasso regression to see which one is more accurate. As we can see, Ridge regression has a lower RMSE which is more accurate for our model, so we will plot the Cross validation curve. We get the RSME score from the Metrics Package.

### Plotting Ridge regression CV plot

```
plot(cv_ride)
abline(v=log(best_ride),col="red",lty=2)
```



The cross validation ridge plot is shown above. According to “<https://bookdown.org/ssjackson300/Machine-Learning-Lecture-Notes/choosing-lambda.html>”:

“What is plotted is the estimated CV MSE for each value of (log)lambda on the x-axis. The dotted line on the far left indicates the value of lambda which minimizes CV error. The dotted line roughly in the middle of the x-axis indicates the 1-standard-error lambda- recall that this is the maximum value that lambda can take while still falling within the on standard error interval of the minimum-CV lambda. The second line of code has manually added a dot-dash horizontal line at the upper end of the 1-standard deviation interval of the MSE at the minimum-CV lambda to illustrate this point further”. These plots can change with randomization according to our seed number.

The summary shown by the trained/tested data are regularized and explain the scale of the variables within the ridge regression. We can use the test dataframe metrics to find the ideal candidate for the Atlanta Falcons on the offensive side of the ball.


### Plotting our Comparison graph between MLR and Ridge Regression MLR

```
plot(test_df$pred_ols, test_df$TD,
     xlim = range(c(test_df$pred_ols, test_df$pred_ride)),
     ylim = range(c(test_df$pred_ols, test_df$pred_ride)),
```

```

xlab="Predicted TD", ylab="Actual TD",
main="OLS (blue) vs Ridge (red)", pch=19, col=rgb(0,0,1,0.6))
points(test_df$pred_ride, test_df$TD, pch=17, col=rgb(1,0,0,0.6))
abline(0,1,lty=2)
legend("topleft", legend=c("OLS","Ridge"), pch=c(19,17),
      col=c("blue","red"), bty="n")

```



We can compare our MLR Ordinary Least Squares Regression Model with our Cross-Validated, Ridge Regression Model visually as shown above.

### Summary of trained and tested data metrics

```
summary(train_df)
```

##	Rk	Player	From	To
##	Min. : 1.0	Length:143	Min. :1992	Min. :1992
##	1st Qu.: 79.0	Class :character	1st Qu.:2001	1st Qu.:2004
##	Median :156.0	Mode :character	Median :2009	Median :2012
##	Mean :166.7		Mean :2009	Mean :2011
##	3rd Qu.:263.0		3rd Qu.:2018	3rd Qu.:2020
##	Max. :326.0		Max. :2024	Max. :2024
##	G	Pos	AV	Tgt
##	Min. : 1.0	Length:143	Min. : 0.00	Min. : 1.00
##	1st Qu.: 12.0	Class :character	1st Qu.: 1.00	1st Qu.: 4.50
##	Median : 22.0	Mode :character	Median : 3.00	Median : 28.00
##	Mean : 35.8		Mean : 11.43	Mean : 82.94
##	3rd Qu.: 52.5		3rd Qu.: 11.00	3rd Qu.: 96.50
##	Max. :222.0		Max. :203.00	Max. :1320.00
##	Rec	Ctch%	Yds	Y/R

```
## Min. : 1.00 Min. :0.1430 Min. : -7.0 Min. : -7.000
## 1st Qu.: 2.00 1st Qu.:0.5240 1st Qu.: 19.5 1st Qu.: 7.100
## Median : 16.00 Median :0.6670 Median : 181.0 Median :10.000
## Mean : 53.17 Mean :0.6708 Mean : 618.0 Mean : 9.727
## 3rd Qu.: 60.50 3rd Qu.:0.7830 3rd Qu.: 622.0 3rd Qu.:13.300
## Max. :848.00 Max. :1.0000 Max. :12896.0 Max. :21.000
## TD Lng Y/Tgt R/G
## Min. : 0.00 Min. : -5.0 Min. : -7.000 Min. :0.000
## 1st Qu.: 0.00 1st Qu.:12.5 1st Qu.: 4.700 1st Qu.:0.250
## Median : 1.00 Median :28.0 Median : 6.300 Median :0.800
## Mean : 3.65 Mean :33.3 Mean : 5.987 Mean :1.262
## 3rd Qu.: 3.00 3rd Qu.:53.0 3rd Qu.: 8.000 3rd Qu.:1.800
## Max. :60.00 Max. :94.0 Max. :14.500 Max. :6.300
## Y/G Fmb
## Min. : -0.80 Min. : 0.0
## 1st Qu.: 1.90 1st Qu.: 0.0
## Median : 8.40 Median : 0.0
## Mean :14.51 Mean : 3.0
## 3rd Qu.:19.65 3rd Qu.: 2.5
## Max. :95.50 Max. :89.0
```

```
summary(test_df)
```

```
## Rk Player From To
## Min. : 2.0 Length:36 Min. :1992 Min. :1993
## 1st Qu.: 83.5 Class :character 1st Qu.:1999 1st Qu.:2000
## Median :159.0 Mode :character Median :2009 Median :2014
## Mean :154.6 Mean :2008 Mean :2010
## 3rd Qu.:226.8 3rd Qu.:2018 3rd Qu.:2019
## Max. :301.0 Max. :2023 Max. :2024
## G Pos AV Tgt
## Min. : 2.00 Length:36 Min. : 0.00 Min. : 1.00
## 1st Qu.: 15.75 Class :character 1st Qu.: 1.00 1st Qu.: 8.75
## Median : 30.50 Mode :character Median : 2.00 Median : 26.00
## Mean : 39.97 Mean : 10.53 Mean : 124.97
## 3rd Qu.: 46.25 3rd Qu.: 11.00 3rd Qu.: 92.50
## Max. :171.00 Max. :107.00 Max. :1377.00
## Rec Ctch% Yds Y/R
## Min. : 1.00 Min. :0.2500 Min. : 6.00 Min. : 2.00
## 1st Qu.: 5.75 1st Qu.:0.5560 1st Qu.: 44.75 1st Qu.: 7.70
## Median : 15.50 Median :0.6410 Median : 175.00 Median : 9.35
## Mean : 76.19 Mean :0.6502 Mean : 903.67 Mean :10.76
## 3rd Qu.: 56.50 3rd Qu.:0.7255 3rd Qu.: 547.75 3rd Qu.:12.20
## Max. :808.00 Max. :1.0000 Max. :10863.00 Max. :35.00
## TD Lng Y/Tgt R/G
## Min. : 0.000 Min. : 8.00 Min. : 1.200 Min. :0.000
## 1st Qu.: 0.000 1st Qu.:16.50 1st Qu.: 4.775 1st Qu.:0.300
## Median : 1.000 Median :25.50 Median : 6.100 Median :0.600
## Mean : 5.472 Mean :34.19 Mean : 6.958 Mean :1.264
## 3rd Qu.: 3.250 3rd Qu.:50.00 3rd Qu.: 7.900 3rd Qu.:1.925
## Max. :63.000 Max. :90.00 Max. :26.000 Max. :4.700
## Y/G Fmb pred_ols pred_ridge
## Min. : 0.400 Min. : 0.000 Min. : -0.00366 Min. : 0.2363
## 1st Qu.: 3.025 1st Qu.: 0.000 1st Qu.: 0.31546 1st Qu.: 0.5358
```

```
## Median : 6.000   Median : 0.500   Median : 1.09513   Median : 1.2570
## Mean   :13.419   Mean    : 1.778    Mean    : 5.44531   Mean    : 5.2606
## 3rd Qu.:19.600   3rd Qu.: 2.250    3rd Qu.: 3.74536   3rd Qu.: 3.5922
## Max.   :63.500   Max.     :14.000    Max.    :62.03274   Max.    :56.6947
## pred_lasso
## Min.    : 0.5532
## 1st Qu.: 0.8057
## Median : 1.4760
## Mean    : 5.1934
## 3rd Qu.: 3.7754
## Max.    :53.5819
```

## Conclusion of Findings

We can now safely say that the Atlanta Falcons TDs can be predicted by multiple factors within a game such as catch percentage, receptions, yards, and other numerical factors.

We see that an ideal candidate for the Atlanta Falcons on the offensive side of the ball (particularly WRs, TEs and RBs) would have the stats of:

Around 20-30 Receiving targets (Based on Median estimate) (Tgt)

Approximately 41 Receptions (Based on IQR) (Rec)

64% catch percentage (based on Median) (Ctch%)

Total yards of 503 - 903\*\* in a given year (IQR and mean) (Yds)

\*\*note that the mean is not used as a measure of spread here, but rather a range of indication for players with the IQR fitting the offensive scheme of the falcons.

We can assume that a combination of these stats (with slight variability based on outliers with superstar potential) will lead to a productive increase (or stability in case of outliers) in Touchdowns for the Atlanta Falcons in the case of picking up free agents, resigning players, or trading for talent.

Keep in mind that these stats are based on my personal interpretation and can vary from person to person. I have used data online and interpreted the Falcons offensive scheme from Zac Robinson's (Falcons Offensive Coordinator) Air-Raid philosophy (based on what I've found online).

## Future improvements to the model

1. Automation of roster data in future findings
2. A classification model detailing other external factors (behavior, team chemistry, etc.) can also be used in tandem with this model in order to make an even more accurate decision.
3. Expanding the model to look at more advanced offensive stats/metrics like Y/G, Y/Tgt, etc.
4. Making an extensive ML regression workflow to determine team-fit with Free Agent data, NFL Trade data, or College NCAA data for drafts (NCAA would have to be adjusted to NFL standards for accurate comparison specifically for the Falcons).
5. Create multiple models and compare test statistics to figure out which results are more tangible to use based on directions from team scouts, front offices, coaches, etc.

## References Used:

<https://bookdown.org/ssjackson300/Machine-Learning-Lecture-Notes/choosing-lambda.html>

<https://www.pro-football-reference.com/teams/atl/career-receiving.htm>

<https://online.stat.psu.edu/stat462/node/180/>

<https://stats.stackexchange.com/questions/279300/how-to-interpret-cross-validation-plot-from-glmnet>

<https://www.datacamp.com/tutorial/tutorial-lasso-ridge-regression>

<https://online.stat.psu.edu/stat462/node/131/>

[https://www.reddit.com/r/AskStatistics/comments/ycjoy4/what\\_threshold\\_is\\_used\\_to\\_assess\\_multicollinearity/](https://www.reddit.com/r/AskStatistics/comments/ycjoy4/what_threshold_is_used_to_assess_multicollinearity/)

<https://www.datacamp.com/doc/r/regression>

<https://www.geo.fu-berlin.de/en/v/soga-r/Basics-of-statistics/Linear-Regression/Polynomial-Regression/Polynomial-Regression---An-example/index.html>

<https://online.stat.psu.edu/stat462/node/177/>

[https://scikit-learn.org/stable/modules/generated/sklearn.linear\\_model.Ridge.html](https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.Ridge.html)

[https://www.reddit.com/r/rstats/comments/oibm2w/how\\_to\\_deal\\_with\\_multicollinearity\\_in\\_linear/](https://www.reddit.com/r/rstats/comments/oibm2w/how_to_deal_with_multicollinearity_in_linear/)

<https://www.datacamp.com/tutorial/multicollinearity>

<https://www.rdocumentation.org/packages/car/versions/3.1-3>

<https://www.rdocumentation.org/packages/Metrics/versions/0.1.4>

<https://www.r-bloggers.com/2021/10/lambda-min-lambda-1se-and-cross-validation-in-lasso-binomial-response/>

<https://www.techtarget.com/searchenterpriseai/definition/data-splitting#:~:text=Training%20sets%20are%20commonly%20used,the%20final%20model%20works%20correctly.>

<https://www.datacamp.com/doc/r/scatterplot-in-r>

<https://www.educba.com/multiple-linear-regression-in-r/>

<https://www.si.com/nfl/falcons/inside-zac-robinson-s-falcons-offense-speed-shifts-controlled-chaos-01hxvw1caw08?>